




shortcut: identity, or stride-2 subsample and zero-pad — no convolution, so never hooked



the filter narrows: $\sigma = 1.00 \rightarrow 0.30$ over epochs 0–20
then $\sigma = 0$ from epoch 21: an exact identity, no kernel applied

early7: only these 7 sites are filtered



conv 3×3, no bias

16 / 32 / 64 channels at stages 1 / 2 / 3, on maps of 32×32 / 16×16 / 8×8



G_σ insertion, 19 of them

separable, depthwise, reflection-padded; one fixed 9×9 support (radius 4) for every σ
that support spans 28% / 56% / 112% of the map width at stages 1 / 2 / 3



one residual block, drawn above

two hooked convolutions and one shortcut; the numbers beneath give its two site indices

insertion mask — **all19**: every site is filtered; **early7**: only sites 0–6 are, and the 12 deeper sites run unfiltered

both names are campaign configuration tokens, as in R32_Gplateau_input_bilinear_all19