



A Where the resolution is reduced

Each arm uses exactly one gate. Both are real spatial changes: nothing is upsampled back, and $r = 32$ is an exact bypass. Drawn at $r = 16$, the low leg of Rprog.




shrinks the float image before $\mathcal{N}$, so every site – the stem included – sees 16×16



keeps the input at 32 and shrinks the stem output instead, so the stem still sees 32×32 and only the blocks are reduced

B The four resolution schedules, over 30 epochs

Darker means a smaller map. Every arm is back at 32×32 from epoch 12, well before training ends.



C The two reduction operators

*bilinear is an antialiased resample; *max is adaptive max pooling, whose windows are not always disjoint.

*bilinear: `F.interpolate(..., mode="bilinear", align_corners=False, antialias=True)`

*max: `F.adaptive_max_pool2d`

32 → 16: windows all of width 2, 0 overlapping pairs – disjoint.

32 → 24: windows all of width 2, but 16 consecutive pairs *overlap*:

