



Analysis B: resolution-only checkpoint after epoch 6 (last updates at $r=16$), seed 0, absolute CE, recalibrated BatchNorm per state.
 Weights, centre, directions (pair 0), subsets and axes are identical across states; only the block-1 reduction changes.
 Counterfactual evaluations of one set of weights, not three trained models.

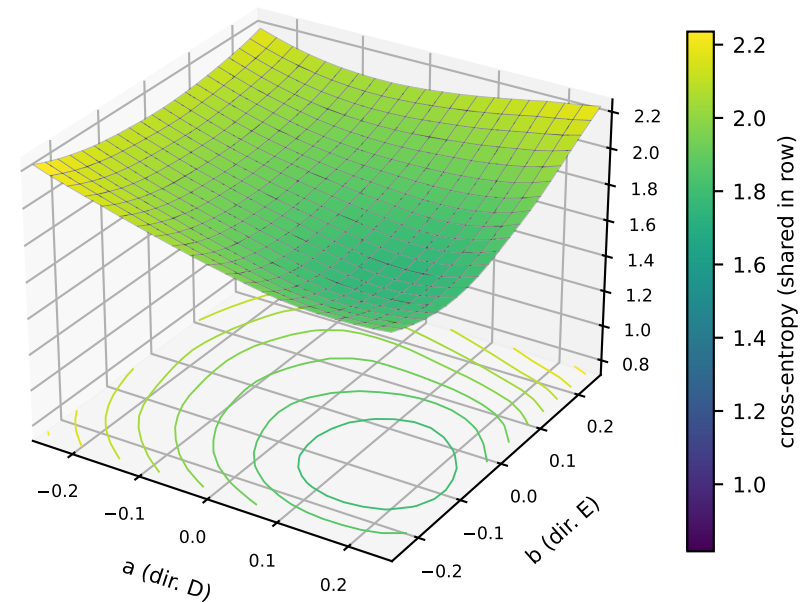
evaluated at $r=16$
 train probe (1,000 imgs), centre CE 0.817



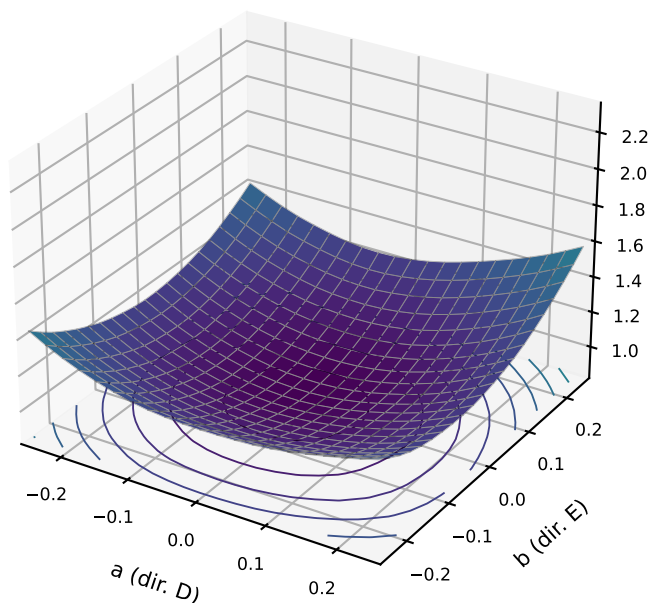
evaluated at $r=24$
 train probe (1,000 imgs), centre CE 1.065



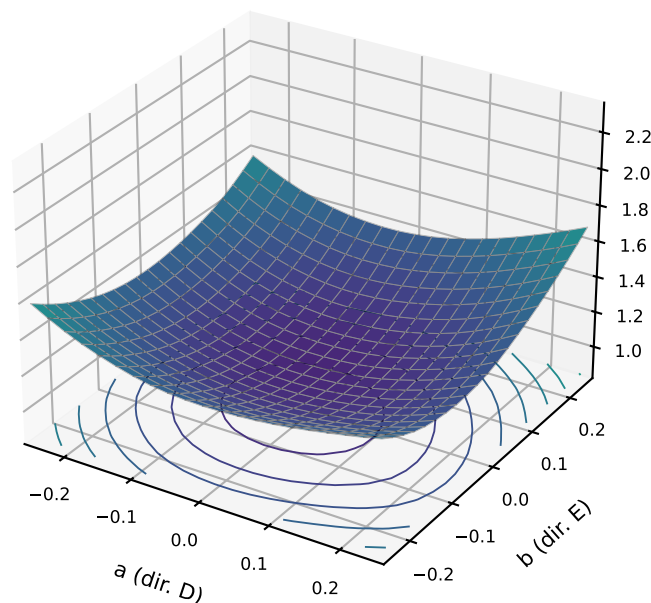
evaluated at $r=32$ (target)
 train probe (1,000 imgs), centre CE 1.847



evaluated at $r=16$
 test probe (1,000 imgs), centre CE 0.921



evaluated at $r=24$
 test probe (1,000 imgs), centre CE 1.073



evaluated at $r=32$ (target)
 test probe (1,000 imgs), centre CE 1.925

