



Analysis B: resolution-only checkpoint after epoch 6 (last updates at  $r=16$ ), seed 0, centre-subtracted CE, recalibrated BatchNorm per state.  
 Weights, centre, directions (pair 0), subsets and axes are identical across states; only the block-1 reduction changes.  
 Counterfactual evaluations of one set of weights, not three trained models.


evaluated at  $r=16$   
 train probe (1,000 imgs), centre CE 0.817




evaluated at  $r=24$   
 train probe (1,000 imgs), centre CE 1.065




evaluated at  $r=32$  (target)  
 train probe (1,000 imgs), centre CE 1.847



evaluated at  $r=16$   
 test probe (1,000 imgs), centre CE 0.921



evaluated at  $r=24$   
 test probe (1,000 imgs), centre CE 1.073



evaluated at  $r=32$  (target)  
 test probe (1,000 imgs), centre CE 1.925

